

Visualization

```
In [2]: import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import poisson
```

Descriptive Statistics

```
In [4]: df=pd.read_csv('breast_cancer.csv')
df.describe()
```

```
Out[4]:
```

	clump_thickness	cell_size_uniformity	cell_shape_uniformity	marginal_adhesion	single_epithelia
count	683.000000	683.000000	683.000000	683.000000	6
mean	4.442167	3.150805	3.215227	2.830161	
std	2.820761	3.065145	2.988581	2.864562	
min	1.000000	1.000000	1.000000	1.000000	
25%	2.000000	1.000000	1.000000	1.000000	
50%	4.000000	1.000000	1.000000	1.000000	
75%	6.000000	5.000000	5.000000	4.000000	
max	10.000000	10.000000	10.000000	10.000000	


Correlation Heatmap

```
In [5]: plt.figure(figsize=(16, 6))

# define the mask to set the values in the upper triangle to True
mask = np.triu(np.ones_like(df.corr(), dtype=bool))

heatmap = sns.heatmap(df.corr(), vmin=-1, vmax=1, mask=mask, annot=True, cmap='gist_rainbow')

heatmap.set_title('Correlation Heatmap');
```



In [2]:

```
import numpy as np
import pandas as pd
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
```

In [3]:

```
df=pd.read_csv(r'C:\Users\ajsru\OneDrive\Desktop\SCM516\breast_cancer.csv')
```

In [4]:

```
df.head(5)
```

Out[4]:

	clump_thickness	cell_size_uniformity	cell_shape_uniformity	marginal_adhesion	single_epithelial_cell_size	bare_nuclei	blank
0	5	1	1	1	2	1	
1	5	4	4	5	7	10	
2	3	1	1	1	2	2	
3	6	8	8	1	3	4	
4	4	1	1	3	2	1	

In [5]:

```
y=df.loc[:, 'diagnose']
```

In [6]:

```
X=df.drop('diagnose', axis=1)
```

In [7]:

```
Xtrain, Xtest, Ytrain, Ytest=train_test_split(X,y, test_size=0.3, random_state=0)
```

In [8]:

```
Xtrain;
```

In [9]:

```
type(Xtrain)
```

Out[9]:

```
pandas.core.frame.DataFrame
```

In [10]:

```
np.shape(Xtrain)
```

Out[10]:

```
(478, 9)
```

In [11]:

```
model=GaussianNB()
```

In [12]:

```
model.fit(Xtrain,Ytrain)
```

Out[12]:

▼ GaussianNB

GaussianNB()


In [13]:

```
Xtest
```

Out[13]:

	clump_thickness	cell_size_uniformity	cell_shape_uniformity	marginal_adhesion	single_epithelial_cell_size	bare_nuclei	blastocyst
113	1	1	1	1	2	5	0
378	3	1	1	1	2	1	0
303	5	5	5	2	5	10	0
504	4	7	8	3	4	10	0
301	1	1	1	1	2	1	0
...
521	5	1	1	1	2	1	0
647	1	1	3	1	2	1	0
503	3	1	2	2	2	1	0
498	3	1	1	1	1	1	0
293	10	8	4	4	4	10	0

205 rows × 9 columns



In [14]:

```
ypred=model.predict(Xtest)
```

In [15]:

```
ypred
```

Out[15]:

```
array([0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0,
       1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0,
       0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0,
       0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1,
       1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0,
       1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0,
       1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0,
       0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0,
       0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       1, 0, 0, 0, 0, 0, 0, 1], dtype=int64)
```

In [16]:

```
Ytest
```

Out[16]:

```
113    0
378    0
303    1
504    1
301    0
...
521    0
```

```
521      0
647      0
503      0
498      0
293      1
Name: diagnose, Length: 205, dtype: int64
```

In [17]:

```
model.score(Xtrain,Ytrain)
```

Out[17]:

```
0.9707112970711297
```

In [18]:

```
model.score(Xtest,Ytest)
```

Out[18]:

```
0.9463414634146341
```

In [19]:

```
cm=confusion_matrix(ypred, Ytest)
cm
```

Out[19]:

```
array([[121,  2],
       [ 9,  73]], dtype=int64)
```

In [20]:

```
print(classification_report(ypred, Ytest))
```


	precision	recall	f1-score	support
0	0.93	0.98	0.96	123
1	0.97	0.89	0.93	82
accuracy			0.95	205
macro avg	0.95	0.94	0.94	205
weighted avg	0.95	0.95	0.95	205

In [21]:

```
sns.heatmap(cm, annot=True)
```

Out[21]:

<AxesSubplot:>



In [22]:

```
import pandas as pd
```

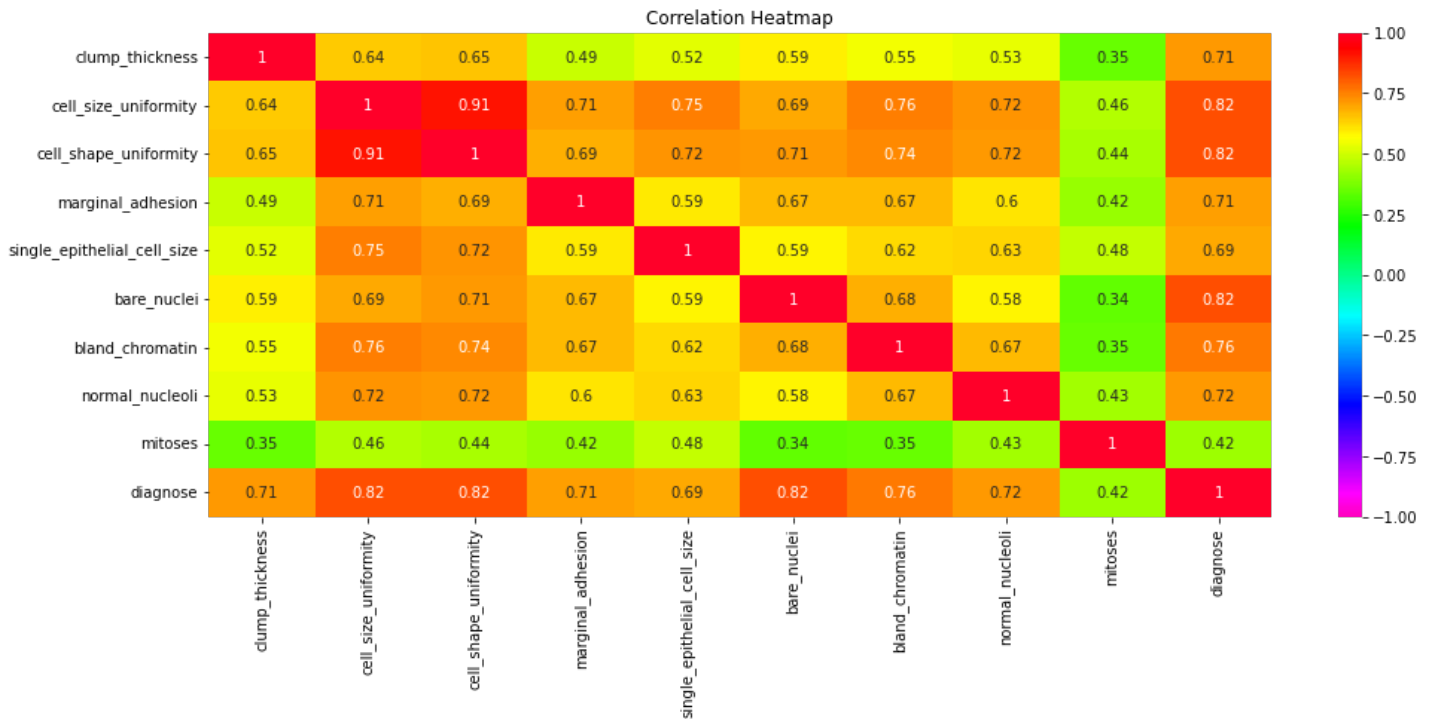
```
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
```

In [23]:

```
plt.figure(figsize=(16, 6))

heatmap = sns.heatmap(df.corr(), vmin=-1, vmax=1, annot=True, cmap='gist_rainbow_r')

heatmap.set_title('Correlation Heatmap');
```




In [24]:

```
sns.lmplot(data=df, x='clump_thickness', y='cell_size_uniformity', hue='diagnose')
```

Out[24]:

<seaborn.axisgrid.FacetGrid at 0x1ed51e2a9a0>

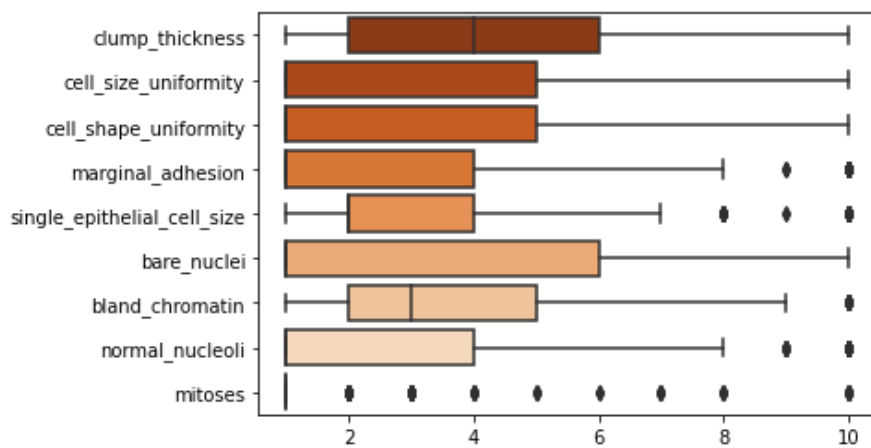


In [25]:

```
sns.boxplot(data=X, orient='h', palette='Oranges_r')
```

Out[25]:

<AxesSubplot:>

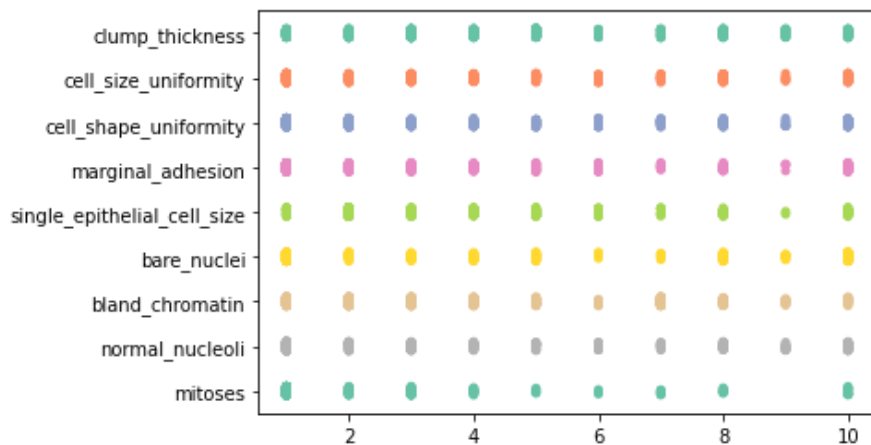


In [26]:

```
sns.stripplot(data=X, orient='h', palette='Set2')
```

Out[26]:

<AxesSubplot:>




In [27]:

```
sns.lmplot(data=df, x='clump_thickness', y='cell_size_uniformity', palette='Set3')
```

Out[27]:

<seaborn.axisgrid.FacetGrid at 0x1ed51fb4f10>




In [28]:

```
sns.lmplot(data=df, x='clump_thickness', y='cell_shape_uniformity')
```

Out[28]:

<seaborn.axisgrid.FacetGrid at 0x1ed523a7850>




In [29]:

```
sns.lmplot(data=df, x='clump_thickness', y='normal_nucleoli')
```

Out[29]:

<seaborn.axisgrid.FacetGrid at 0x1ed51506eb0>



In [30]:

```
sns.heatmap(cm, annot=True, cmap='RdPu');
```

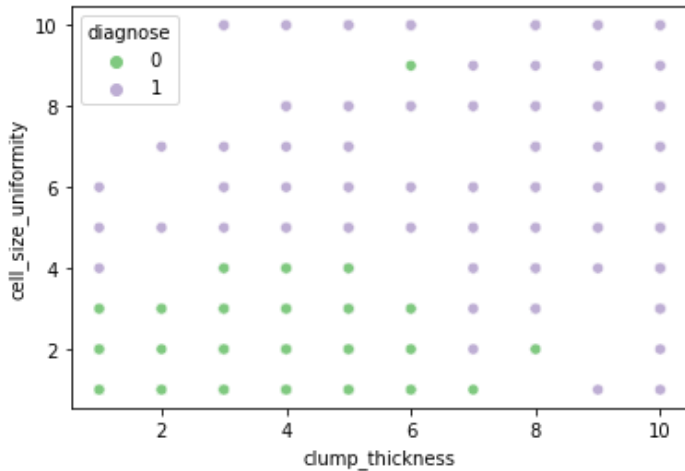




In [35]:

```
import seaborn as sns

sns.scatterplot(x=X.iloc[:,0], y=X.iloc[:,1], hue=y, palette='Accent');
```



In [32]:

```
!pip install xgboost
```

Requirement already satisfied: xgboost in c:\users\ajsru\anaconda3\lib\site-packages (1.6.2)
Requirement already satisfied: numpy in c:\users\ajsru\anaconda3\lib\site-packages (from xgboost) (1.21.5)
Requirement already satisfied: scipy in c:\users\ajsru\anaconda3\lib\site-packages (from xgboost) (1.7.3)


In [33]:

```
from sklearn.metrics import confusion_matrix, classification_report, plot_roc_curve, roc_curve, auc
```

In [34]:

```
disp=plot_roc_curve(model, Xtest, Ytest)
plt.plot([0,1], [0,1], color='orange', linestyle='--');
```

C:\Users\ajsru\AppData\Roaming\Python\Python39\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.
warnings.warn(msg, category=FutureWarning)



KNN Classification Model

Creating Training and Testing Models

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
from sklearn import neighbors, datasets
from matplotlib.colors import ListedColormap
from sklearn.inspection import DecisionBoundaryDisplay
```

```
In [3]: df=pd.read_csv('breast_cancer.csv')
```

```
In [4]: df.head(5)
```

```
Out[4]:
```

	clump_thickness	cell_size_uniformity	cell_shape_uniformity	marginal_adhesion	single_epithelial_cel
0	5	1	1	1	
1	5	4	4	5	
2	3	1	1	1	
3	6	8	8	1	
4	4	1	1	3	

```
In [5]: X=df.drop('diagnose', axis=1)
```

```
In [6]: X1=df.iloc[:,0:4]
```

```
In [7]: X1.head()
```

```
Out[7]:
```

	clump_thickness	cell_size_uniformity	cell_shape_uniformity	marginal_adhesion
0	5	1	1	1
1	5	4	4	5
2	3	1	1	1
3	6	8	8	1
4	4	1	1	3

```
In [8]: xlabel=X1.columns[0]
ylabel=X1.columns[1]
```

```
In [9]: y=df['diagnose']
```

```
In [10]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=
```

```
In [11]: # Data standardization
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
# Fit only on X_train
```

```
scaler.fit(X_train)
```

```
# Scale both X_train and X_test
```

```
X_train = scaler.transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

```
In [12]: model=KNeighborsClassifier(n_neighbors=5) # k=5
```

```
In [13]: model.fit(X_train, y_train)
```

```
Out[13]: ▼ KNeighborsClassifier
KNeighborsClassifier()
```

```
In [14]: ypred=model.predict(X_test)
```

```
In [15]: ypred
```

```
Out[15]: array([0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0,
        1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0,
        0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0,
        0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1,
        1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0,
        1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0,
        1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0,
        0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0,
        0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        1, 0, 0, 0, 0, 0, 1], dtype=int64)
```

Accuracy for Training Data

```
In [16]: model.score(X_train, y_train)
```

```
Out[16]: 0.9790794979079498
```

Accuracy for Testing Data


```
In [17]: model.score(X_test, y_test)
```

```
Out[17]: 0.9512195121951219
```

Confusion Matrix

```
In [18]: cm=confusion_matrix(ypred, y_test)
```

```
In [19]: sns.heatmap(cm, annot=True, cmap='RdPu');
```



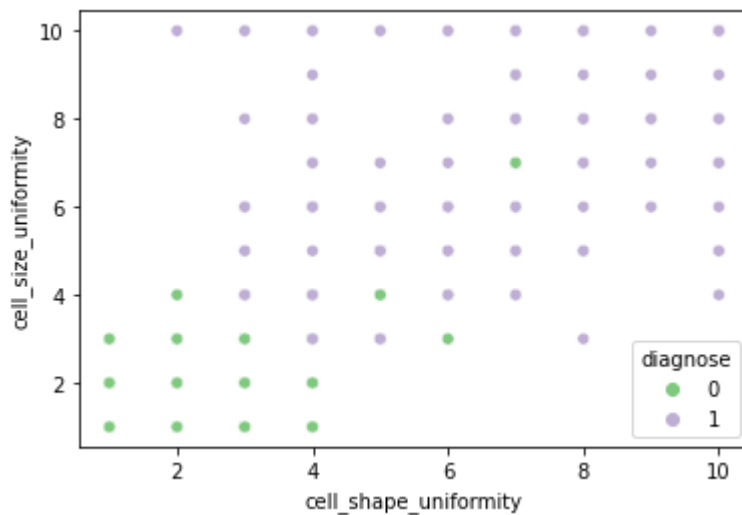
Classification Report

```
In [20]: print(classification_report(ypred, y_test))
```

	precision	recall	f1-score	support
0	0.97	0.95	0.96	132
1	0.92	0.95	0.93	73
accuracy			0.95	205
macro avg	0.94	0.95	0.95	205
weighted avg	0.95	0.95	0.95	205

Scatter Plot - Cell_shape_uniformity vs Cell_size_uniformity

```
In [21]: sns.scatterplot(data=df, x="cell_shape_uniformity", y="cell_size_uniformity", hue="dis")
```



Plotting decision boundry when considering two features (Clump Thickness and Cell size uniformity)

```
In [22]: cmap_light = ListedColormap(["orange", "cornflowerblue"])
cmap_bold = ["darkorange", "darkblue"]
```

```
In [23]: n_neighbors = 5

X=df.iloc[:,0:2]
y=df['diagnose']

cmap_light = ListedColormap(["orange", "cornflowerblue"])
cmap_bold = ["darkorange", "darkblue"]

for weights in ["uniform", "distance"]:
    # we create an instance of Neighbours Classifier and fit the data.
    model = neighbors.KNeighborsClassifier(n_neighbors, weights=weights)
    model.fit(X, y)

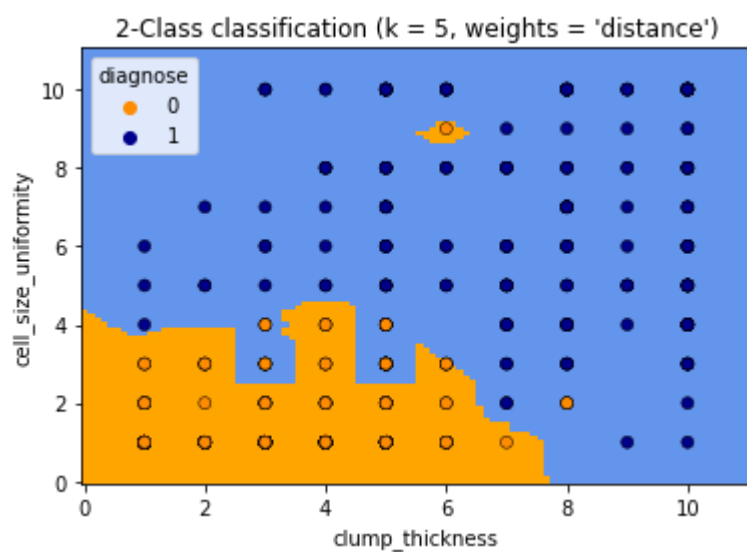
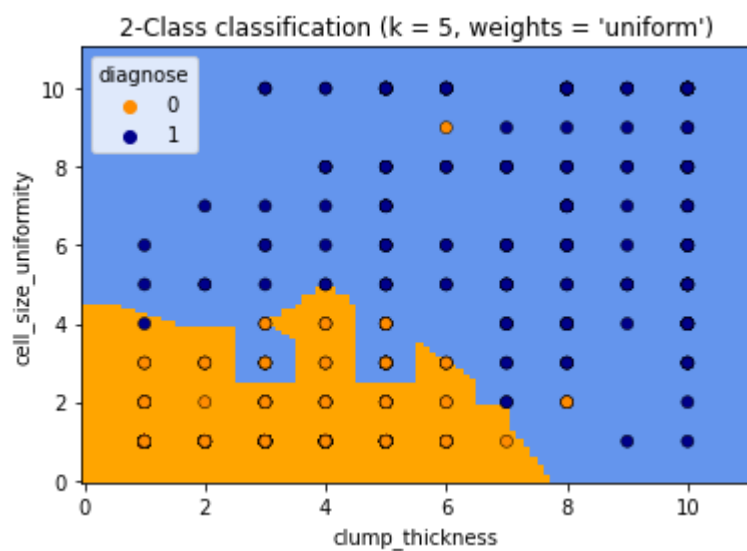
    _, ax= plt.subplots()
    DecisionBoundaryDisplay.from_estimator(model,X,cmap=cmap_light,ax=ax,response_method='predict',
        xlabel=X.columns[0],
        ylabel=X.columns[1],
        shading="auto")

    # Plot also the training points'

    sns.scatterplot(
        x=X.iloc[:,0],
        y=X.iloc[:,1],
        hue=y,
        palette=cmap_bold,
        alpha=1.0,
        edgecolor="black",
    )

    plt.title(
```

```
"2-Class classification (k = %i, weights = '%s')" % (n_neighbors, weights));
```



Random Forest

```
In [30]: import numpy as np
import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import *
import seaborn as sns
```

```
In [11]: df = pd.read_csv(r'C:\Users\91735\Downloads\breast_cancer.csv')
```

```
In [12]: df.describe()
```

```
Out[12]:
```

	clump_thickness	cell_size_uniformity	cell_shape_uniformity	marginal_adhesion	single_epithelia
count	683.000000	683.000000	683.000000	683.000000	6
mean	4.442167	3.150805	3.215227	2.830161	
std	2.820761	3.065145	2.988581	2.864562	
min	1.000000	1.000000	1.000000	1.000000	
25%	2.000000	1.000000	1.000000	1.000000	
50%	4.000000	1.000000	1.000000	1.000000	
75%	6.000000	5.000000	5.000000	4.000000	
max	10.000000	10.000000	10.000000	10.000000	

```
In [3]: x = df.drop('diagnose',axis = 1)
y = df['diagnose']
```

```
In [4]: xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size = 0.3,random_state=0)
```

```
In [18]: model = RandomForestClassifier(random_state = 1,n_estimators = 2000)
```

```
In [19]: model.fit(x,y)
```

```
Out[19]:
```

RandomForestClassifier
RandomForestClassifier(n_estimators=2000, random_state=1)

```
In [21]: I = model.feature_importances_
```

```
In [22]: I
```

```
Out[22]: array([0.04786191, 0.27018697, 0.22232477, 0.02821389, 0.08151435,
0.16604124, 0.10599434, 0.07083451, 0.00702803])
```

```
In [ ]:
```



```
In [6]: model.fit(xtrain,ytrain)
        model.score(xtrain,ytrain)
```

```
Out[6]: 1.0
```

```
In [7]: model.score(xtest,ytest)
```

```
Out[7]: 0.9512195121951219
```


```
In [8]: ypred = model.predict(xtest)
```

```
In [28]: cm = confusion_matrix(ypred,ytest)
        cm
```

```
Out[28]: array([[126,  6],
               [ 4, 69]], dtype=int64)
```

```
In [31]: sns.heatmap(cm, annot = True,  cmap = 'RdPu')
```

```
Out[31]: <AxesSubplot:>
```



```
In [10]: print(classification_report(ypred,ytest))
```

	precision	recall	f1-score	support
0	0.97	0.95	0.96	132
1	0.92	0.95	0.93	73
accuracy			0.95	205
macro avg	0.94	0.95	0.95	205
weighted avg	0.95	0.95	0.95	205

Logistics Regression

```
In [20]: # Importing Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
import seaborn as sns
from mlxtend.plotting import plot_decision_regions
from sklearn.decomposition import PCA
```

```
In [21]: # Loading dataset
df = pd.read_csv(r'breast_cancer.csv')

df.describe()
```

```
Out[21]:
```

	clump_thickness	cell_size_uniformity	cell_shape_uniformity	marginal_adhesion	single_epithelia
count	683.000000	683.000000	683.000000	683.000000	6
mean	4.442167	3.150805	3.215227	2.830161	
std	2.820761	3.065145	2.988581	2.864562	
min	1.000000	1.000000	1.000000	1.000000	
25%	2.000000	1.000000	1.000000	1.000000	
50%	4.000000	1.000000	1.000000	1.000000	
75%	6.000000	5.000000	5.000000	4.000000	
max	10.000000	10.000000	10.000000	10.000000	

```
In [22]: # Previewing data real quick
df.head(10)
```

Out[22]:

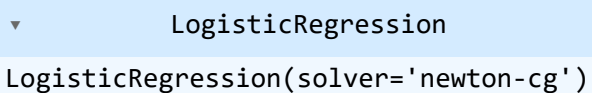
	clump_thickness	cell_size_uniformity	cell_shape_uniformity	marginal_adhesion	single_epithelial_cel
0	5	1	1	1	
1	5	4	4	5	
2	3	1	1	1	
3	6	8	8	1	
4	4	1	1	3	
5	8	10	10	8	
6	1	1	1	1	
7	2	1	2	1	
8	2	1	1	1	
9	4	2	1	1	

In [23]: *# Defining x variables and y variable*
 x = df.iloc[:,0:9].values
 y = df.iloc[:,9].values

In [24]: *# Splitting data for training and testing*
 xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size = 0.3, random_state =

In [25]: *# Initializing the model*
 lrmodel = LogisticRegression(solver='newton-cg')

In [26]: *# Fitting our training data on the model*
 lrmodel.fit(xtrain, ytrain)

Out[26]:  LogisticRegression(solver='newton-cg')

In [27]: *# Checking intercept of our regression model*
 lrmodel.intercept_

Out[27]: array([-9.17796322])

In [28]: *# Checking the coefficients of our regression model*
 lrmodel.coef_

Out[28]: array([[0.3943735 , 0.14430425, 0.21099364, 0.19743802, 0.22348085,
 0.46479598, 0.31518682, 0.24472243, 0.18561027]])

In [29]: *# Predicting our testing data*
 lrprediction = lrmodel.predict(xtest)

In [30]: *# Checking accuracy of our model by comparing it with training data*
 trscore = lrmodel.score(xtrain, ytrain)

```
In [31]: # Checking accuracy of our model by comparing it with testing data
lr_score = lr_model.score(x_test, y_test)
```


```
In [32]: # Our model has an accuracy of 95.6 %
print('Training Accuracy = ', tr_score * 100, '%')
print('Testing Accuracy = ', lr_score * 100, '%')
```

```
Training Accuracy = 97.48953974895397 %
Testing Accuracy = 95.60975609756098 %
```

```
In [33]: confusion_matrix = confusion_matrix(lr_prediction, y_test)
```

```
In [34]: sns.heatmap(confusion_matrix, annot = True)
```

```
Out[34]: <AxesSubplot:>
```



```
In [35]: print(classification_report(lr_prediction, y_test))
```

	precision	recall	f1-score	support
0	0.97	0.96	0.97	131
1	0.93	0.95	0.94	74
accuracy			0.96	205
macro avg	0.95	0.95	0.95	205
weighted avg	0.96	0.96	0.96	205

```
In [36]: df
```

Out[36]:

	clump_thickness	cell_size_uniformity	cell_shape_uniformity	marginal_adhesion	single_epithelial_c
0	5	1	1	1	
1	5	4	4	5	
2	3	1	1	1	
3	6	8	8	1	
4	4	1	1	3	
...
678	3	1	1	1	
679	2	1	1	1	
680	5	10	10	3	
681	4	8	6	4	
682	4	8	8	5	

683 rows × 10 columns

In [40]:

```
# Creating Decision Boundaries between features with high coefficient (clump thickness)

dbx = df.iloc[:,[0,5]].values
dby = df['diagnose'].values

dbxtrain, dbxtest, dbytrain, dbytest = train_test_split(dbx, dby, test_size=0.3, random_state=0)


lr = LogisticRegression(solver='newton-cg', random_state=0)

lr.fit(dbxtrain, dbytrain);

fig8, ax = plt.subplots(figsize=(10, 10))
fig8 = plot_decision_regions(dbxtrain, dbytrain, clf=lr, legend=2);

plt.show
```

Out[40]: <function matplotlib.pyplot.show(close=None, block=None)>



ROC curve

In [2]: `!pip install xgboost`

Requirement already satisfied: xgboost in c:\users\kjosep14\anaconda3\lib\site-packages (1.6.2)
 Requirement already satisfied: scipy in c:\users\kjosep14\anaconda3\lib\site-packages (from xgboost) (1.7.3)
 Requirement already satisfied: numpy in c:\users\kjosep14\anaconda3\lib\site-packages (from xgboost) (1.21.5)

In []: `import numpy as np
 import pandas as pd
 import matplotlib.pyplot as plt
 import seaborn as sns
 from sklearn.model_selection import train_test_split
 from sklearn.metrics import confusion_matrix, classification_report, plot_roc_curve, r
 from sklearn.ensemble import RandomForestClassifier
 from sklearn.naive_bayes import GaussianNB
 from sklearn.linear_model import LogisticRegression
 from sklearn.neighbors import KNeighborsClassifier
 from sklearn.tree import DecisionTreeClassifier
 from sklearn.svm import SVC
 from sklearn.ensemble import AdaBoostClassifier
 from sklearn.preprocessing import StandardScaler
 from sklearn.model_selection import train_test_split
 import xgboost as xgb
 import shap`

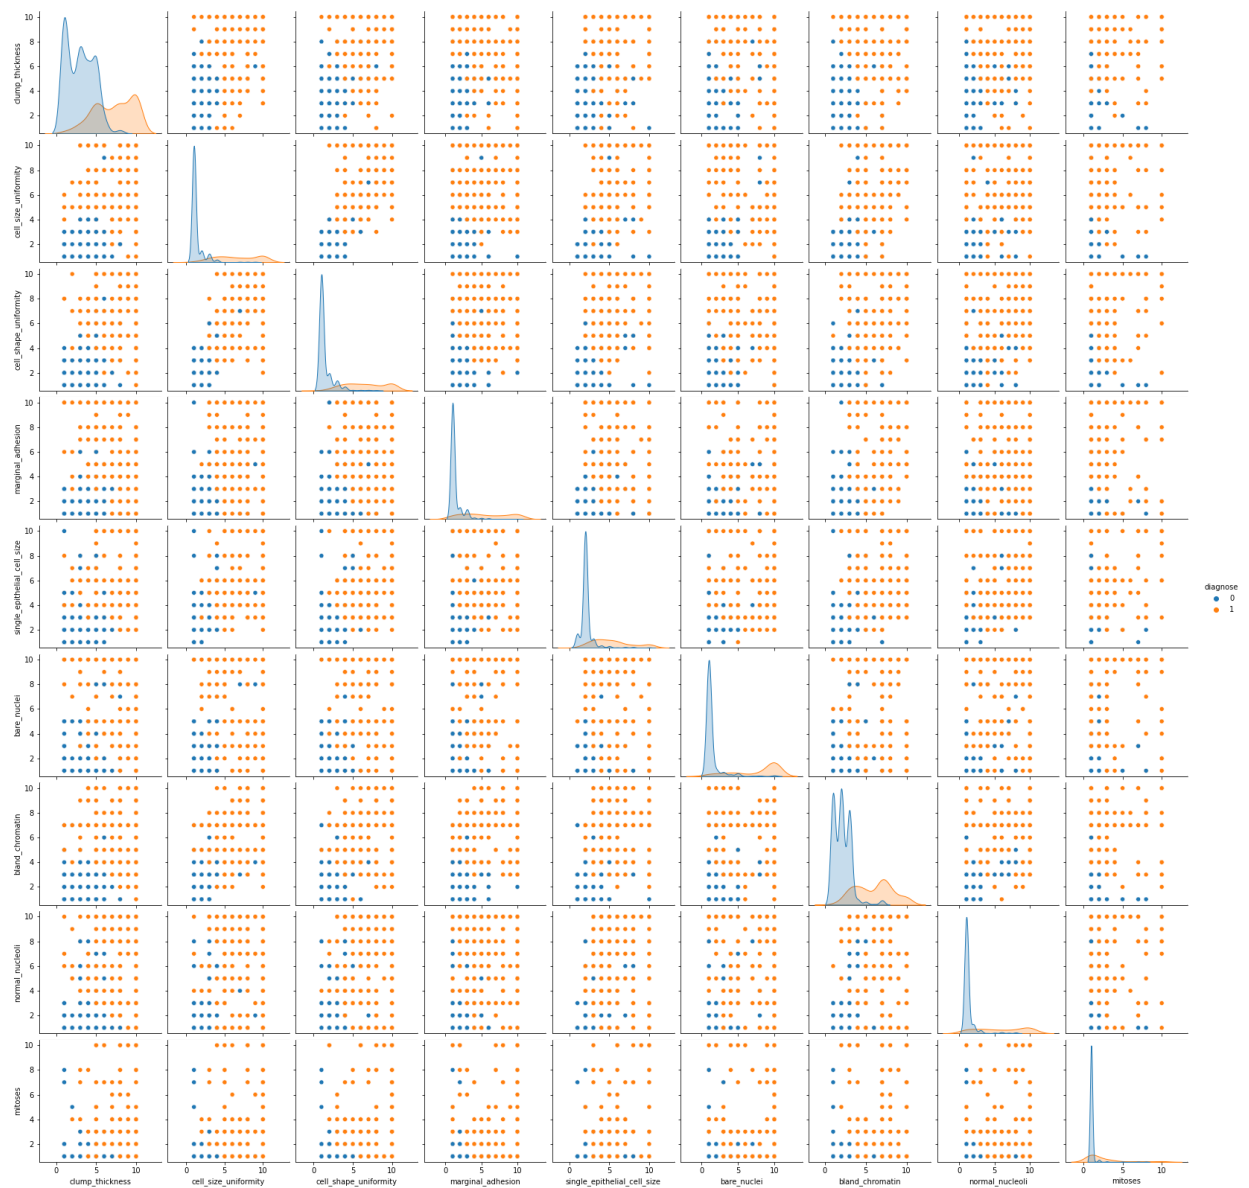
In [6]: `df=pd.read_csv('breast_cancer.csv')`

In [7]: `df.head(3)`

Out[7]:

	clump_thickness	cell_size_uniformity	cell_shape_uniformity	marginal_adhesion	single_epithelial_cel
0	5	1	1	1	
1	5	4	4	5	
2	3	1	1	1	

In [8]: `sns.pairplot(df, hue='diagnose');` *# pairplot of dataset, scatter plot which represents each pair of features*



```
In [9]: X=df.drop('diagnose', axis=1)
        y=df['diagnose']
```

```
In [10]: # split data into testing and training

Xtrain, Xtest, ytrain, ytest=train_test_split(X,y, test_size=0.3, random_state=0)
```

```
In [11]: # Standardaize data

scaler=StandardScaler()
scaler.fit(Xtrain)
Xtrain=scaler.transform(Xtrain)
Xtest=scaler.transform(Xtest)
```

```
In [50]: # define all models

model1 = LogisticRegression(random_state=1,solver='newton-cg',multi_class='multinomial')
model2 = RandomForestClassifier(random_state=1, n_estimators=2000)
model3 = GaussianNB()
model4 = KNeighborsClassifier(n_neighbors=5)
```


In [51]: *# fit the model*

```
model1.fit(Xtrain, ytrain)
model2.fit(Xtrain, ytrain)
model3.fit(Xtrain, ytrain)
model4.fit(Xtrain, ytrain);
```

In [52]: *# predict class (y) for testing data (Xtest)*

```
ypred1=model1.predict(Xtest)
ypred2=model2.predict(Xtest)
ypred3=model3.predict(Xtest)
ypred4=model4.predict(Xtest)
```

In [53]: ypred6;

In [54]:

```
print("Accuracy of Logistic Regression is: ", model1.score(Xtrain, ytrain))
print("Accuracy of Random Forest is: ", model2.score(Xtrain, ytrain))
print("Accuracy of Gaussaian NB is: ", model3.score(Xtrain, ytrain))
print("Accuracy of KNN is: ", model6.score(Xtrain, ytrain))
```


```
Accuracy of Logistic Regression is:  0.9769874476987448
Accuracy of Random Forest is:  1.0
Accuracy of Gaussaian NB is:  0.9707112970711297
Accuracy of KNN is:  0.9790794979079498
```

In [56]: *# calculate confusion matrix*


```
cm1=confusion_matrix(ypred1, ytest)
cm2=confusion_matrix(ypred2, ytest)
cm3=confusion_matrix(ypred3, ytest)
cm4=confusion_matrix(ypred4, ytest)
```

In [57]: *# plot confusion matrix*

```
class_names=[0,1] # name of classes
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)
sns.heatmap(cm1, annot=True, cmap="YlGnBu" ,fmt='g') # cm1 is the confusion matrix j
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion matrix of Logistic Regression model', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label');
```



```
In [58]: # Confusion matrix for model 1 ---- cm1 is the confusion matrix for model 1 (Logistic
group_names = ['True Neg', 'False Pos', 'False Neg', 'True Pos']
group_counts = ["{0:0.0f}".format(value) for value in
                 cm1.flatten()]
group_percentages = ["{0:.2%}".format(value) for value in
                     cm1.flatten()/np.sum(cm1)]
labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in
           zip(group_names, group_counts, group_percentages)]
labels = np.asarray(labels).reshape(2,2)
sns.heatmap(cm1, annot=labels, fmt='', cmap='Blues');
```




```
In [59]: # confusion matrix for model 2 (Random Forest)
group_names = ['True Neg', 'False Pos', 'False Neg', 'True Pos']
group_counts = ["{0:0.0f}".format(value) for value in
                 cm2.flatten()]
group_percentages = ["{0:.2%}".format(value) for value in
                     cm2.flatten()/np.sum(cm2)]
labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in
           zip(group_names, group_counts, group_percentages)]
labels = np.asarray(labels).reshape(2,2)
sns.heatmap(cm2, annot=labels, fmt='', cmap='Blues');
```

```
group_percentages = [{"0:.2%}".format(value) for value in
                     cm2.flatten()/np.sum(cm2)]

labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in
          zip(group_names,group_counts,group_percentages)]

labels = np.asarray(labels).reshape(2,2)
sns.heatmap(cm2, annot=labels, fmt='', cmap='Blues');
```



```
In [46]: # confusion matrix for model 3 (NB)


group_names = ['True Neg', 'False Pos', 'False Neg', 'True Pos']

group_counts = [{"0:0.0f}".format(value) for value in
                cm3.flatten()]

group_percentages = [{"0:.2%}".format(value) for value in
                     cm3.flatten()/np.sum(cm3)]

labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in
          zip(group_names,group_counts,group_percentages)]

labels = np.asarray(labels).reshape(2,2)
sns.heatmap(cm3, annot=labels, fmt='', cmap='Blues');
```



```
In [60]: print(classification_report(ypred1, ytest)) # first model (Logistic Regression)
```

	precision	recall	f1-score	support
0	0.97	0.96	0.97	131
1	0.93	0.95	0.94	74
accuracy			0.96	205
macro avg	0.95	0.95	0.95	205
weighted avg	0.96	0.96	0.96	205

In [61]: *### plot ROC curve for Xtrain and ytrain*

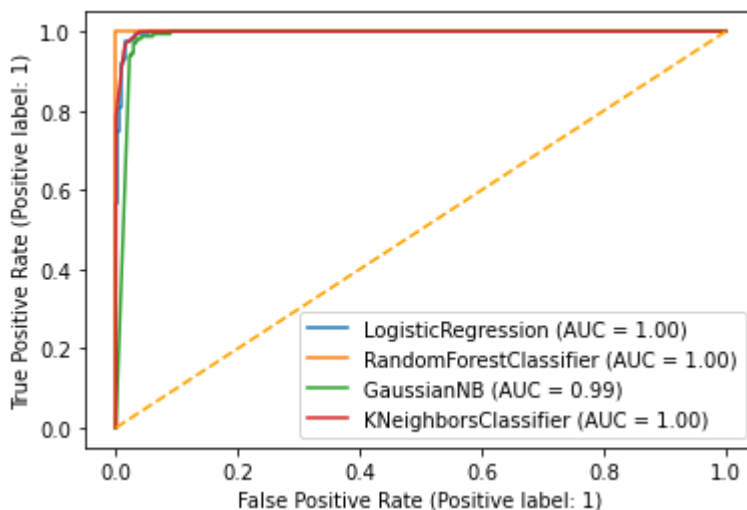
```
disp=plot_roc_curve(model1, Xtrain, ytrain)
plot_roc_curve(model2, Xtrain, ytrain, ax=disp.ax_)
plot_roc_curve(model3, Xtrain, ytrain, ax=disp.ax_)
plot_roc_curve(model4, Xtrain, ytrain, ax=disp.ax_)
plt.plot([0,1], [0,1], color='orange', linestyle='--');
```

Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.



In [62]: *### Plot ROC Curve for Xtest, ytest*


```
disp=plot_roc_curve(model1, Xtest, ytest)
plot_roc_curve(model2, Xtest, ytest, ax=disp.ax_)
plot_roc_curve(model3, Xtest, ytest, ax=disp.ax_)
plot_roc_curve(model4, Xtest, ytest, ax=disp.ax_)
plt.plot([0,1], [0,1], color='orange', linestyle='--');
```

Function `plot_roc_curve` is deprecated; Function `:func:`plot_roc_curve`` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: `:meth:`sklearn.metrics.RocCurveDisplay.from_predictions`` or `:meth:`sklearn.metrics.RocCurveDisplay.from_estimator``.

Function `plot_roc_curve` is deprecated; Function `:func:`plot_roc_curve`` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: `:meth:`sklearn.metrics.RocCurveDisplay.from_predictions`` or `:meth:`sklearn.metrics.RocCurveDisplay.from_estimator``.

Function `plot_roc_curve` is deprecated; Function `:func:`plot_roc_curve`` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: `:meth:`sklearn.metrics.RocCurveDisplay.from_predictions`` or `:meth:`sklearn.metrics.RocCurveDisplay.from_estimator``.

Function `plot_roc_curve` is deprecated; Function `:func:`plot_roc_curve`` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: `:meth:`sklearn.metrics.RocCurveDisplay.from_predictions`` or `:meth:`sklearn.metrics.RocCurveDisplay.from_estimator``.



Feature importance using Random Forest

```
In [64]: rf=RandomForestClassifier(random_state=1, n_estimators=100)
```

```
In [65]: rf.fit(X,y)
```

```
Out[65]: RandomForestClassifier
RandomForestClassifier(random_state=1)
```

```
In [66]: I=rf.feature_importances_
```

```
In [67]: I
```

```
Out[67]: array([0.05398005, 0.2516356 , 0.2464734 , 0.02746694, 0.10317641,
        0.13508699, 0.11311992, 0.06109489, 0.00796582])
```

```
In [68]: X[0:2]
```

Out[68]:

	clump_thickness	cell_size_uniformity	cell_shape_uniformity	marginal_adhesion	single_epithelial_cel
0	5	1	1	1	
1	5	4	4	5	

In [69]:

```
df2=pd.DataFrame({
    "Importance": [0.05398005, 0.2516356 , 0.2464734 , 0.02746694, 0.10317641,
                  0.13508699, 0.11311992, 0.06109489, 0.00796582],
    "diagnose": ["clump_thickness", "cell_size_uniformity", "cell_shape_uniformity", "marginal_adhesion", "single_epithelial_cell_size", "bare_nuclei", "bland_chromatin", "normal_nucleoli", "mitoses"] })
```

In [70]: df2

Out[70]:

	Importance	diagnose
0	0.053980	clump_thickness
1	0.251636	cell_size_uniformity
2	0.246473	cell_shape_uniformity
3	0.027467	marginal_adhesion
4	0.103176	single_epithelial_cell_size
5	0.135087	bare_nuclei
6	0.113120	bland_chromatin
7	0.061095	normal_nucleoli
8	0.007966	mitoses

In [71]: sns.barplot(data=df2, y="diagnose", x="Importance");

